

sklearn.feature_selection.RFECV

```
class sklearn.feature_selection.RFECV(estimator, step=1, min_features_to_select=1, cv=None, scoring=None, verbose=0, n_jobs=None)
```

[\[source\]](#)

Feature ranking with recursive feature elimination and cross-validated selection of the best number of features.

See glossary entry for [cross-validation estimator](#).

Read more in the [User Guide](#).

Parameters:

estimator : object

A supervised learning estimator with a `fit` method that provides information about feature importance either through a `coef_` attribute or through a `feature_importances_` attribute.

step : int or float, optional (default=1)

If greater than or equal to 1, then `step` corresponds to the (integer) number of features to remove at each iteration. If within (0.0, 1.0), then `step` corresponds to the percentage (rounded down) of features to remove at each iteration. Note that the last iteration may remove fewer than `step` features in order to reach `min_features_to_select`.

min_features_to_select : int, (default=1)

The minimum number of features to be selected. This number of features will always be scored, even if the difference between the original feature count and `min_features_to_select` isn't divisible by `step`.

cv : int, cross-validation generator or an iterable, optional

Determines the cross-validation splitting strategy. Possible inputs for `cv` are:

- None, to use the default 5-fold cross-validation,
- integer, to specify the number of folds.
- [CV splitter](#),
- An iterable yielding (train, test) splits as arrays of indices.

For integer/None inputs, if `y` is binary or multiclass, [sklearn.model_selection.StratifiedKFold](#) is used. If the estimator is a classifier or if `y` is neither binary nor multiclass, [sklearn.model_selection.KFold](#) is used.

Refer [User Guide](#) for the various cross-validation strategies that can be used here.

Changed in version 0.22: cv default value of None changed from 3-fold to 5-fold.

scoring : string, callable or None, optional, (default=None)

A string (see model evaluation documentation) or a scorer callable object / function with signature `scorer(estimator, X, y)`.

verbose : int, (default=0)

Controls verbosity of output.

n_jobs : int or None, optional (default=None)

Number of cores to run in parallel while fitting across folds. None means 1 unless in a [joblib.parallel_backend](#) context. -1 means using all processors. See [Glossary](#) for more details.

Attributes:

n_features_ : int

The number of selected features with cross-validation.

support_ : array of shape [n_features]

The mask of selected features.

ranking_ : array of shape [n_features]

The feature ranking, such that `ranking_[i]` corresponds to the ranking position of the *i*-th feature. Selected (i.e., estimated best) features are assigned rank 1.

grid_scores_ : array of shape [n_subsets_of_features]

The cross-validation scores such that `grid_scores_[i]` corresponds to the CV score of the *i*-th subset of features.

Toggle Menu

estimator : object

The external estimator fit on the reduced dataset.

See also:

[RFE](#)

Recursive feature elimination

Notes

The size of `grid_scores_` is equal to `ceil((n_features - min_features_to_select) / step) + 1`, where `step` is the number of features removed at each iteration.

Allows NaN/Inf in the input if the underlying estimator does as well.

References

R6f4d61ceb411-1 Guyon, I., Weston, J., Barnhill, S., & Vapnik, V., "Gene selection for cancer classification using support vector machines", Mach. Learn., 46(1-3), 389–422, 2002.

Examples

The following example shows how to retrieve the a-priori not known 5 informative features in the Friedman #1 dataset.

```
>>> from sklearn.datasets import make_friedman1
>>> from sklearn.feature_selection import RFECV
>>> from sklearn.svm import SVR
>>> X, y = make_friedman1(n_samples=50, n_features=10, random_state=0)
>>> estimator = SVR(kernel="linear")
>>> selector = RFECV(estimator, step=1, cv=5)
>>> selector = selector.fit(X, y)
>>> selector.support_
array([ True,  True,  True,  True,  True, False, False, False, False,
        False])
>>> selector.ranking_
array([1, 1, 1, 1, 1, 6, 4, 3, 2, 5])
```

Methods

decision_function (self, X)	Compute the decision function of x.
fit (self, X, y[, groups])	Fit the RFE model and automatically tune the number of selected
fit_transform (self, X[, y])	Fit to data, then transform it.
get_params (self[, deep])	Get parameters for this estimator.
get_support (self[, indices])	Get a mask, or integer index, of the features selected
inverse_transform (self, X)	Reverse the transformation operation
predict (self, X)	Reduce X to the selected features and then predict using the
predict_log_proba (self, X)	Predict class log-probabilities for X.
predict_proba (self, X)	Predict class probabilities for X.
score (self, X, y)	Reduce X to the selected features and then return the score of the
set_params (self, **params)	Set the parameters of this estimator.
transform (self, X)	Reduce X to the selected features.

```
__init__(self, estimator, step=1, min_features_to_select=1, cv=None, scoring=None, verbose=0, n_jobs=None)
```

[\[source\]](#)

Initialize self. See help(type(self)) for accurate signature.

```
decision_function(self, X)
```

[\[source\]](#)

Compute the decision function of x.

Parameters:

X : {array-like or sparse matrix} of shape (n_samples, n_features)

The input samples. Internally, it will be converted to `dtype=np.float32` and if a sparse matrix is provided to a `csc_matrix`.

Returns:

score : array, shape = [n_samples, n_classes] or [n_samples]

The decision function of the input samples. The order of the classes corresponds to that in the attribute `classes_`. Regression and binary classification produce an array of shape [n_samples].

Fit the RFE model and automatically tune the number of selected features.

Parameters:

X : {array-like, sparse matrix} of shape (n_samples, n_features)

Training vector, where n_samples is the number of samples and n_features is the total number of features.

y : array-like of shape (n_samples,)

Target values (integers for classification, real numbers for regression).

groups : array-like of shape (n_samples,) or None

Group labels for the samples used while splitting the dataset into train/test set. Only used in conjunction with a "Group" [cv](#) instance (e.g., [GroupKFold](#)).

```
fit_transform(self, X, y=None, **fit_params)
```

[\[source\]](#)

Fit to data, then transform it.

Fits transformer to X and y with optional parameters fit_params and returns a transformed version of X.

Parameters:

X : numpy array of shape [n_samples, n_features]

Training set.

y : numpy array of shape [n_samples]

Target values.

****fit_params : dict**

Additional fit parameters.

Returns:

X_new : numpy array of shape [n_samples, n_features_new]

Transformed array.

```
get_params(self, deep=True)
```

[\[source\]](#)

Get parameters for this estimator.

Parameters:

deep : bool, default=True

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

params : mapping of string to any

Parameter names mapped to their values.

```
get_support(self, indices=False)
```

[\[source\]](#)

Get a mask, or integer index, of the features selected

Parameters:

indices : boolean (default False)

If True, the return value will be an array of integers, rather than a boolean mask.

Returns:

support : array

An index that selects the retained features from a feature vector. If indices is False, this is a boolean array of shape [# input features], in which an element is True iff its corresponding feature is selected for retention. If indices is True, this is an integer array of shape [# output features] whose values are indices into the input feature vector.

```
transform(self, X)
```

[\[source\]](#)

Reverse the transformation operation

Parameters:

***X* : array of shape [n_samples, n_selected_features]**

The input samples.

Returns:

***X_r* : array of shape [n_samples, n_original_features]**

x with columns of zeros inserted where features would have been removed by [transform](#).

`predict(self, X)`

[\[source\]](#)

Reduce X to the selected features and then predict using the
underlying estimator.

Parameters:

***X* : array of shape [n_samples, n_features]**

The input samples.

Returns:

***y* : array of shape [n_samples]**

The predicted target values.

`predict_log_proba(self, X)`

[\[source\]](#)

Predict class log-probabilities for X.

Parameters:

***X* : array of shape [n_samples, n_features]**

The input samples.

Returns:

***p* : array of shape (n_samples, n_classes)**

The class log-probabilities of the input samples. The order of the classes corresponds to that in the attribute [classes](#).

`predict_proba(self, X)`

[\[source\]](#)

Predict class probabilities for X.

Parameters:

***X* : {array-like or sparse matrix} of shape (n_samples, n_features)**

The input samples. Internally, it will be converted to dtype=np.float32 and if a sparse matrix is provided to a sparse csr_matrix.

Returns:

***p* : array of shape (n_samples, n_classes)**

The class probabilities of the input samples. The order of the classes corresponds to that in the attribute [classes](#).

`score(self, X, y)`

[\[source\]](#)

Reduce X to the selected features and then return the score of the
underlying estimator.

Parameters:

***X* : array of shape [n_samples, n_features]**

The input samples.

***y* : array of shape [n_samples]**

The target values.

`set_params(self, **params)`

[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

`params : dict`**
Estimator parameters.

Returns:

`self : object`
Estimator instance.

`transform(self, X)`

[\[source\]](#)

Reduce X to the selected features.

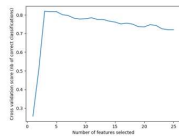
Parameters:

`X : array of shape [n_samples, n_features]`
The input samples.

Returns:

`X_r : array of shape [n_samples, n_selected_features]`
The input samples with only the selected features.

Examples using `sklearn.feature_selection.RFECV`



[Recursive feature
elimination with cross-
validation](#)

© 2007 - 2019, scikit-learn developers (BSD License). [Show this page source](#)